

# Data Mining Homework 3

*Shan Jiang*

*10/27/2019*

## Contents

Problem 2 . . . . .	1
Load packages . . . . .	1
Read in data . . . . .	1
(a) Construct a linear support vector classifier . . . . .	1
Model . . . . .	2
Test error and its standard error . . . . .	3
(b) Construct a support vector classifier with Radial kernel . . . . .	4
Tune parameter . . . . .	4
Test error and its standard error . . . . .	4
Plots . . . . .	5
(c) AdaBoost algorithm-classifier . . . . .	5
Test Error and std.error . . . . .	6
Summary . . . . .	7

## Problem 2

### Load packages

```
library(tidyverse)
library(broom)
library(kernlab)      # SVM methodology
library(e1071)        # SVM methodology
library(glmnet)
require(ggplot2)
```

### Read in data

```
## 250 obs. train data
train_data <- read.delim("./synth.tr", header = T, sep = "")
## 1000 obs. test data
test_data <- read.delim("./synth.te", header = T, sep = "")
```

There are in total in the train dataframe, containing two predictors and a binary outcome.

### (a) Construct a linear support vector classifier

## Model

We must encode the response as a factor variable.

```
train_data = train_data %>%  
  mutate(yc= as.factor(yc))  
test_data = test_data %>%  
  mutate(yc= as.factor(yc))
```

using support vector classifier, in this case, we try to tune the best parameter:

```
## Using scaled data  
set.seed(1)  
tune.out = tune(svm, yc~.,  
  data= train_data,  
  kernel="linear",  
  ranges= list(cost=c(0.1,1,10,100,1000),  
    gamma= c(0.5,1,2,3,4)))  
summary(tune.out)
```

```
##  
## Parameter tuning of 'svm':  
##  
## - sampling method: 10-fold cross validation  
##  
## - best parameters:  
##   cost gamma  
##     1    0.5  
##  
## - best performance: 0.14  
##  
## - Detailed performance results:  
##   cost gamma error dispersion  
## 1  1e-01   0.5 0.152 0.09003703  
## 2  1e+00   0.5 0.140 0.07831560  
## 3  1e+01   0.5 0.140 0.07831560  
## 4  1e+02   0.5 0.140 0.07831560  
## 5  1e+03   0.5 0.140 0.07831560  
## 6  1e-01   1.0 0.152 0.09003703  
## 7  1e+00   1.0 0.140 0.07831560  
## 8  1e+01   1.0 0.140 0.07831560  
## 9  1e+02   1.0 0.140 0.07831560  
## 10 1e+03   1.0 0.140 0.07831560  
## 11 1e-01   2.0 0.152 0.09003703  
## 12 1e+00   2.0 0.140 0.07831560  
## 13 1e+01   2.0 0.140 0.07831560  
## 14 1e+02   2.0 0.140 0.07831560  
## 15 1e+03   2.0 0.140 0.07831560  
## 16 1e-01   3.0 0.152 0.09003703  
## 17 1e+00   3.0 0.140 0.07831560  
## 18 1e+01   3.0 0.140 0.07831560  
## 19 1e+02   3.0 0.140 0.07831560  
## 20 1e+03   3.0 0.140 0.07831560  
## 21 1e-01   4.0 0.152 0.09003703
```

```
## 22 1e+00    4.0 0.140 0.07831560
## 23 1e+01    4.0 0.140 0.07831560
## 24 1e+02    4.0 0.140 0.07831560
## 25 1e+03    4.0 0.140 0.07831560
```

```
svmfit1 = tune.out$best.model
#summary(svmfit1)
```

We see that `cost = 1` results in the lowest cross-validation error rate. The `tune()` function stores the best model obtained.

### Test error and its standard error

```
ypred1 = predict(svmfit1, test_data, type = "response")
test.error1 <- 1 - sum(ypred1 == test_data$yc) / nrow(test_data)
test.error1
```

```
## [1] 0.107
```

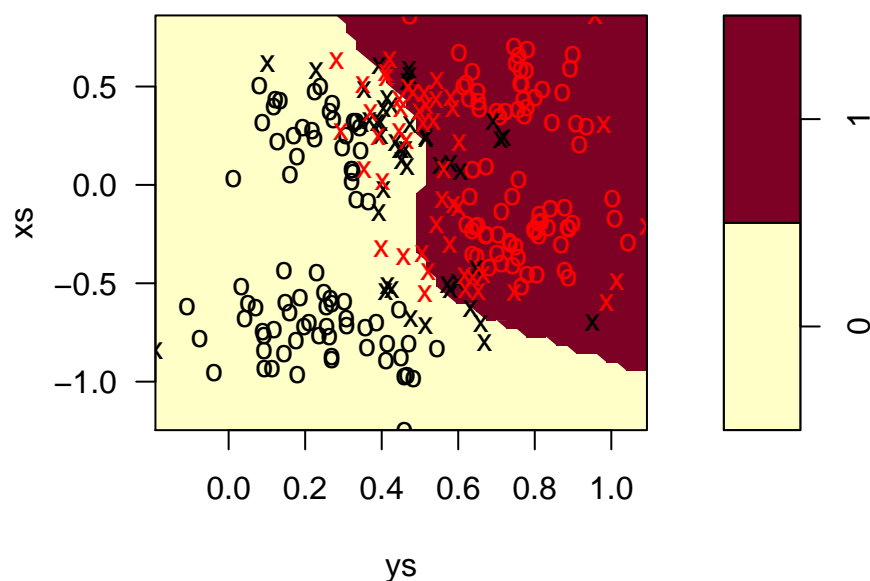
```
std.error1 <- sqrt((test.error1 * (1 - test.error1)) / (nrow(test_data)))
std.error1
```

```
## [1] 0.009775019
```

```
## use the best-fit parameter to fit the model
svmfit1_model = svm(yc ~ ., data = train_data,
                    kernel = "radial", gamma = 0.5, cost = 1)

# summary(svmfit1_model)
plot(svmfit1_model, train_data)
```

## SVM classification plot



(b) Construct a support vector classifier with Radial kernel

### Tune parameter

Perform cross-validation using `tune()` to select the best choice.

```
set.seed(1)
tune.out2 = tune(svm, ys~., data=train_data,
                 kernel="radial",
                 ranges= list(cost=c(0.1,1,10,100,1000),
                              gamma= c(0.5,1,2,3,4)))
#summary(tune.out2)
```

Therefore, the best choice of parameters involves `cost=1` and `gamma=0.5`.

```
## use the best-fit parameter to fit the model
svmfit2 = svm(ys~., data = train_data,
              kernel="radial", gamma=0.5, cost =1)
#summary(svmfit2)
```

This tells us, for instance, that a linear kernel was used with `cost = 1`, and that there were 95 support vectors, 47 in one class and 48 in the other.

### Test error and its standard error

```
ypred2 = predict(svmfit2,
                 newdata = test_data, type = "response")

ypred2 = predict(svmfit2, test_data, type = "response")

test.error2 <- 1 - sum(ypred2 == test_data$yc) / nrow(test_data)
test.error2
```

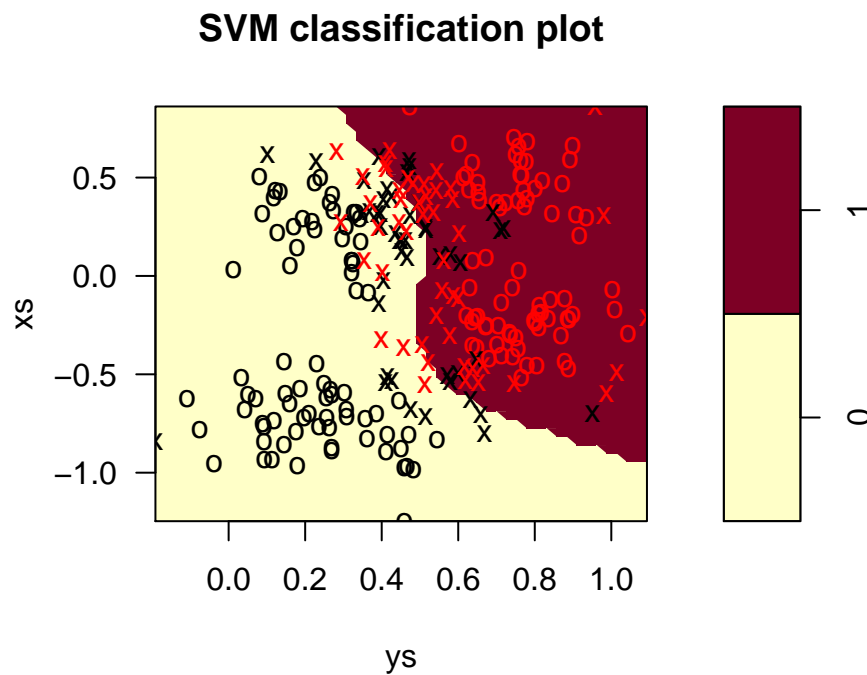
```
## [1] 0.095
```

```
std.error2 <- sqrt((test.error2 * (1 - test.error2)) / (nrow(test_data)))
std.error2
```

```
## [1] 0.00927227
```

Plots

```
plot(svmfit2, train_data)
```



### (c) AdaBoost algorithm-classifier

Boosting is another technique which assumes a weak classifier technique at the beginning, while it attains the true classification.

```
#install.packages("gbm")
library(ada)
```

```
## Loading required package: rpart
```

```
## use of gradient/generalized boosting models in "gbm".
set.seed(1118)
fit.adaboost <- ada(yc ~., data=train_data,
                    iter=50, nu = 0.1, bag.frac=0.5)
summary(fit.adaboost)
```

```
## Call:
## ada(yc ~ ., data = train_data, iter = 50, nu = 0.1, bag.frac = 0.5)
##
## Loss: exponential Method: discrete   Iteration: 50
##
## Training Results
##
## Accuracy: 0.912 Kappa: 0.824
```

Test Error and std.error

```
ypred3 = predict(fit.adaboost,newdata = test_data)

test.error3 <- 1- sum(ypred3 == test_data$yc)/nrow(test_data)
test.error3
```

```
## [1] 0.098
```

```
std.error3 <- sqrt((test.error3 * (1- test.error3))/(nrow(test_data)))
std.error3
```

```
## [1] 0.009401915
```

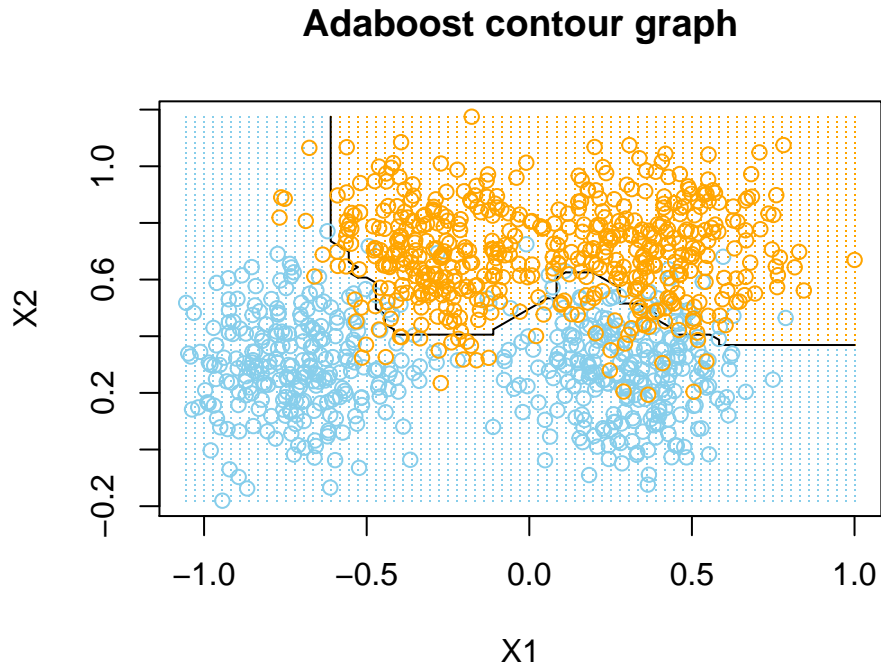
```
## Plotting for adaboost
plot2 = function(model,df_train){
  train = df_train[,1:2]
  L=75
  X=seq(min(train[,1]),max(train[,1]),length=L)
  Y=seq(min(train[,2]),max(train[,2]),length=L)
  XY=expand.grid(X,Y) %>% rename(xs=Var1,ys=Var2)
  yTrain = df_train$yc
  yhat=predict(model,XY)
  colors <- c("SkyBlue", "Orange")
  yhat1 <- colors[as.numeric(yhat)]
  yTrain <- colors[as.numeric(yTrain)]
  plot(train, xlab="X1", ylab="X2",
        xlim = range(train[,1]),
        ylim = range(train[,2]), type="n")
  points(XY,col=yhat1, pch=15,cex=0.1)
```

```

contour(X, Y, matrix(as.numeric(yhat),L,L),
        levels=c(1,2), add=TRUE, drawlabels=FALSE)
points(train, col=yTrain)
title("Adaboost contour graph")
}

plot2(fit.adaboost, test_data)

```



## Summary

```

test.error = c(test.error1, test.error2, test.error3)
sd = c(std.error1, std.error2, std.error3)
method = c("linear support vector classifier",
           "Radial Kernel support vector classifier",
           "Adaboost alogorithm")
summary = data.frame(method,test.error, sd)
knitr::kable(summary)

```

method	test.error	sd
linear support vector classifier	0.107	0.0097750
Radial Kernel support vector classifier	0.095	0.0092723
Adaboost alogorithm	0.098	0.0094019

Comment:

- Three models have slightly different classification test error results, while they vary only a little bit, the linear SVM is of less good fit compared with SVM machine with Kernel model, while Adaboost algorithm perform best.
- The test error and standard deviation are lowest for Adaboost classifier, but we need to be cautious of the overfitting risk.